

PRAVEEN KUMAR S

Full Stack Developer • Python & Django • React.js • AI Systems

+91 73582 57338 | praveensankar3333@gmail.com | Chennai, Tamil Nadu | [GitHub](#) | [LinkedIn](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Full Stack Developer with hands-on experience shipping production-grade web applications using Python, Django REST Framework, React.js. Designed and deployed an offline-first AI assistant running local LLM inference and a cloud-integrated AI code intelligence platform powered by Gemini 2.5 Flash — both built end-to-end, solo. Strong foundation in RESTful API design, component-driven frontend architecture, and CI/CD pipelines via Vercel. Systems-thinking mindset with rapid prototyping, and delivering working software under constraints.

TECHNICAL SKILLS

Languages: Python, JavaScript (ES6+)

Frontend: React.js, HTML5, CSS3, Tailwind CSS, Component Architecture, Hooks, State Management

Backend: Django, Django REST Framework, RESTful API Design, Session Management, Django Allauth

Databases: PostgreSQL, SQLite, Relational Schema Design, ORM (Django)

AI & LLM: Ollama (Qwen, Phi, Mistral, Llama), Local LLM Inference, Google Gemini API, Prompt Engineering

Tools & DevOps: Git, GitHub, CI/CD (Vercel), VS Code, Environment Configuration, Deployment

Core Concepts: REST API Design, MVC/MVT Architecture, SDLC, Data Structures & Algorithms

PROJECTS

Privacy-First Local LLM Orchestration Engine Feb 2026

Stack: Django REST Framework · React.js · Ollama · Qwen 4B / Phi · Python · SQLite

- ▶ Architected a fully offline, zero-cloud-dependency AI assistant running Qwen 4B locally via Ollama — achieving \$0 API cost and 100% data privacy, with average inference response latency under 2500ms on consumer-grade hardware.
- ▶ Built custom Django REST API endpoints managing LLM message routing, inference orchestration, and session state persistence, eliminating reliance on any third-party orchestration framework (no LangChain, no external middleware).
- ▶ Shipped a ChatGPT-style React.js frontend with persistent multi-session memory, real-time streaming UI, and instant model switching across 4 LLMs (Qwen, Phi, Mistral, Llama) — zero page reloads throughout.
- ▶ Tested the engine across 4 LLM families via Ollama to validate inference compatibility and response quality, demonstrating practical understanding of local model behaviour.

GitHub: github.com/PraveenSankar03/local-llm-assistant

AI-Powered Coding Intelligence Studio Jan 2026

Stack: React.js · Django REST Framework · Python · Google Gemini 2.5 Flash API · JavaScript

- ▶ Built a tri-mode code intelligence platform (GENERATE / EXPLAIN / DEBUG) powered by Gemini 2.5 Flash via a Django REST backend — delivering structured AI responses in under 5 seconds end-to-end across all three modes.
- ▶ Consolidated three distinct developer workflows into a single-page React app with zero-reload mode switching, eliminating 100% of context-switching overhead between separate coding tools.
- ▶ Validated prompt stability and response consistency across 5 programming languages (Python, JavaScript, C++, Java, PHP), achieving reliable structured LLM output across diverse language paradigms.
- ▶ Architected a stateless Django REST API layer that fully decouples the AI model from the frontend - enabling future model replacement (GPT-4, Claude) with zero frontend changes; clean separation of concerns by design.

GitHub: github.com/PraveenSankar03/AI-code-explainer-debugger

Full-Stack Social Microblogging Platform Nov 2025

Stack: *Django · Python · SQLite · MVT Architecture · Django Allauth · HTML5 · CSS3*

- ▶ Engineered a full-stack microblogging platform in Django MVT with Django Allauth managing the complete authentication lifecycle (registration, login, session, logout) entirely under framework-native control and no custom auth hacks.
- ▶ Designed normalized relational schemas with ownership-enforced access control: follow/unfollow modeled as a self-referential M2M structure; post authors are the exclusive controllers of their own content via view-layer permission checks.
- ▶ Implemented server-rendered Django views covering post feeds, user profiles, follow graphs, and content management - demonstrating deep command of URL routing, template inheritance, and context processors.
- ▶ Enforced object-level permission checks at the view layer to prevent unauthorized content manipulation, a production-quality security pattern directly applicable to SaaS backends.

GitHub: github.com/PraveenSankar03/django-social-media

ADDITIONAL DEPLOYED PROJECTS

- ▶ Movie Browser App — React.js + TMDB API; live search, dynamic routing, and responsive card layout. Deployed on Vercel. [\[Live Demo\]](#)
- ▶ Weather Fetcher App — React.js + Open-Meteo API; real-time geolocation-based weather with zero-backend architecture. Deployed on Vercel. [\[Live Demo\]](#)

EDUCATION

B.Tech in Biomedical Engineering 2021 – 2025

Dr. M.G.R. Educational and Research Institute, Chennai, Tamil Nadu

CGPA: 8.4 / 10

CERTIFICATIONS & COURSES

The Ultimate Full Stack Development Bootcamp | 2025

Instructors: Kane Ezki, Kalob Taulien · Full-stack web development: JavaScript, Python, Django, React.js

Data Structures and Algorithms (DSA) in Python | 2025

Instructor: Scott Barrett · Big-O complexity, arrays, linked lists, trees, graphs, sorting, dynamic programming

COMPETITIVE PROGRAMMING

LeetCode — Active profile, solving problems and growing

Focusing on arrays, strings, and hash maps, heaps - core patterns used in real-world backend and API development.

CURRENTLY LEARNING

LangChain · Retrieval-Augmented Generation (RAG) · FAISS Vector Search

Extending AI expertise beyond local inference into production-grade LLM pipelines: document retrieval, semantic search with FAISS vector stores, and context-aware generation using LangChain - directly applicable to RAG-based SaaS products and enterprise AI tooling.