

PRAVEEN KUMAR S

Full Stack Developer · Python · Django · React.js · AI-Integrated Systems

+91 73582 57338 | praveensankar3333@gmail.com | Chennai, Tamil Nadu | [GitHub](#) | [LinkedIn](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Software Developer who independently builds production-grade applications end-to-end. Designed and built AI-integrated applications: An offline LLM orchestration engine and Gemini-powered coding intelligence platform, both built solo. Writes clean, decoupled backend APIs and responsive React frontends and knows how to structure systems.

TECHNICAL SKILLS

Languages	Python, JavaScript (ES6+)
Frontend	React.js, Tailwind CSS, HTML5, CSS3, React Hooks, State Management
Backend	Django, Django REST Framework, RESTful APIs
Databases	PostgreSQL, SQLite, Django ORM
AI & LLM	Ollama (Qwen, Phi, Mistral, Llama), Google Gemini API, Prompt Engineering
DevOps	Git, GitHub, CI/CD via Vercel, Deployment
Foundations	MVC/MVT Architecture, Data Structures & Algorithms

PROJECTS / EXPERIENCE

Privacy-First Local LLM Orchestration Engine

Jan 2026

[Django REST Framework](#) · [React.js](#) · [Ollama](#) · [Qwen 4B](#) · [Python](#) · [SQLite](#)

- **Eliminated API cost entirely and maximized data privacy** - Fully offline on consumer hardware via Ollama, sub-7000ms inference latency, zero cloud dependency and no data exposure – ensuring 100% data privacy for the user.
- **Built without LangChain or any third-party orchestration framework** - custom Django REST layer handles all message routing, LLM inference orchestration, and session state management from scratch.
- **Shipped a ChatGPT-grade React frontend** with persistent multi-session memory, real-time streaming UI, and zero page reloads. Tested the engine across 4 LLM families via Ollama to validate inference compatibility and response quality.

➤ github.com/PraveenSankar03/local-llm-assistant

AI-Powered Coding Intelligence Studio

Dec 2025

[React.js](#) · [Django REST Framework](#) · [Google Gemini 2.5 Flash](#) · [Python](#)

- **Unified three developer tools into one platform** – GENERATE (writes code from Scratch) / EXPLAIN (explains the user given code) / DEBUG (debugs and pinpoints errors) in a Single Page React Application, producing end-to-end responses in under 5 seconds.
- **Architected for replaceability** - stateless Django REST layer fully decouples and separates the AI model from the React frontend; swapping from Gemini to GPT-4 or Claude can be done faster with zero frontend changes.
- **Validated across many programming languages** (Python, JavaScript, C++, Java, PHP), confirming structured prompt pipelines for stable and consistent LLM output consistency.

➤ github.com/PraveenSankar03/AI-code-explainer-debugger

Full-Stack Social Microblogging Platform

Oct 2025

[Django](#) · [Python](#) · [SQLite](#) · [MVT Architecture](#) · [Django Allauth](#) · [HTML5](#) · [CSS3](#)

- **Production-quality auth, no custom hacks** - full authentication lifecycle (registration, login, session, logout) managed via Django Allauth entirely under framework-native control – without any third-party framework.
- **Object-level permission enforcement at the view layer** - requires signup to create an account and requires login to post a content. Unauthorized content manipulation blocked at architecture level, directly applicable to SaaS backends.
- **Relational schema designed for scale** - follow/unfollow modeled as a self-referential M2M structure; normalized ownership enforced throughout. Also added feature of dynamic posts rendering – the feed only consists of the posts posted by the users they follow.

➤ github.com/PraveenSankar03/django-social-media

EDUCATION

B.Tech in Biomedical Engineering

CGPA: 8.4 / 10

Dr. M.G.R. Educational and Research Institute, Chennai · 2021 – 2025

CERTIFICATIONS

The Ultimate Full Stack Development Bootcamp | Kane Ezki, Kalob Taulien (2025) · JavaScript, Python, Django, React.js

Data Structures & Algorithms in Python | Scott Barrett (2025) · Big-O, stacks, queues, arrays, trees, graphs, dynamic programming